

Project Report
on
“Real Time Data Streaming Application”

CPSC 531-03 22470

Advanced Database Management

Fall, 2022

Under Guidance Of

Prof. Tseng-Ching James Shen

Department of Computer Science



California State University

Fullerton CA – 92831

December 2022

Prepared By:-

Onkar Muttemwar(885199950) - onkar.muttemwar@csu.fullerton.edu

Sambhaji Ippar(885865899) - sambhaji@csu.fullerton.edu

Aman Rathore(885186841) - aman.r@csu.fullerton.edu

Contents

Introduction	3
Functionalities	3
Architecture Overview	4
Technologies and tools used	5-6
Project skills needed but not limited to	6
Dataset	6-7
GitHub Location of Code	7
Deployment Instructions	8-15
Steps to Run the Application	16
Test Results	16-17
References	18

Introduction


- Real-time Data streaming is a process by which big volumes of data are processed as soon as they are generated as continuous streams.
- According to reports, more than a quarter of the data created would be real-time. There are a lot of data sources that create such type of data, which are IoT sensors, smart devices, and gaming applications which produce data at high volumes and high velocity.
- So there is a need to process this data in real-time as there is some business where real-time processing and analytics are of crucial importance to get an edge over competitors, and also it enables faster decision making along with various other advantages.
- Detecting fraud in real-time, ride-share apps, and e-commerce apps are very important examples of real-time processing.

Functionalities

- Producers are simulating the data stream rapidly, which is then getting ingested by AWS Kinesis in real-time.
- The ingested data in AWS Kinesis is partitioned using shards and sent to Firehose.
- The two connected AWS Firehose ingested the data from Kinesis. One of them is continuously loading the real-time data to Elasticsearch or OpenSearch while the other one is loading the raw data into our S3 data bucket.
- The AWS Elasticsearch or OpenSearch is monitoring our data and producing visual insights using Kibana
- The raw data within the S3 bucket is getting transformed by glue crawlers and moved to Athena
- The AWS Athena is then querying the data and producing the visual insights in Tableau which is connected to Athena server.

Architecture Overview:

Visual Paradigm Online Free Edition



Visual Paradigm Online Free Edition

- The data Stream is a python program which is simulating the data in real-time
- The output data stream is ingested by DataStream which is then processed by KinesisFirehose
- It is then used by OpenSearch for real-time analysis using Kibana
- The other Firehose is used by S3 for Batch processing which is then used to add data in glue tables and analysis using Athena and Tableau

Technologies and tools used:

1. Python



2. Tableau



3. AWS Kinesis



4. Amazon Kinesis



5. AWS S3



6. AWS Glue



7. AWS Athena



8. AWS ElasticSearch



Project skills needed but not limited to:

To work on the project, one must have the following skills but not limited to.

- Having a basic understanding of cloud computing.
- Experience with the AWS platform and its various services.
- Knowledge of python programming or, as an alternative, creating a producer code using JAVA.
- Understanding of security best practices and how they apply to the cloud.

Dataset:

- The architecture of the application is such that it would work on any real-time data set. But for the project's scope, we have used the Bank Marketing Dataset from Kaggle.
- The dataset has a lot of columns from which we can extract a lot of insights that would be extracted to analyze in real-time.
- The dataset can be downloaded from Kaggle

<https://www.kaggle.com/datasets/janiobachmann/bank-marketing-dataset>

1	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	outcome	deposit	
2	59	admin.	married	secondary	no	2343	yes	no	unknown	5	may	1042	1	-1	0	unknown	yes	
3	56	admin.	married	secondary	no	45	no	no	unknown	5	may	1467	1	-1	0	unknown	yes	
4	41	technician	married	secondary	no	1270	yes	no	unknown	5	may	1389	1	-1	0	unknown	yes	
5	55	services	married	secondary	no	2476	yes	no	unknown	5	may	579	1	-1	0	unknown	yes	
6	54	admin.	married	tertiary	no	184	no	no	unknown	5	may	673	2	-1	0	unknown	yes	
7	42	manager	single	tertiary	no	0	yes	yes	unknown	5	may	562	2	-1	0	unknown	yes	
8	56	manager	married	tertiary	no	830	yes	yes	unknown	6	may	1201	1	-1	0	unknown	yes	
9	60	retired	divorced	secondary	no	545	yes	no	unknown	6	may	1030	1	-1	0	unknown	yes	
10	37	technician	married	secondary	no	1	yes	no	unknown	6	may	608	1	-1	0	unknown	yes	
11	28	services	single	secondary	no	5090	yes	no	unknown	6	may	1297	3	-1	0	unknown	yes	
12	38	admin.	single	secondary	no	100	yes	no	unknown	7	may	786	1	-1	0	unknown	yes	
13	30	blue-collar	married	secondary	no	309	yes	no	unknown	7	may	1574	2	-1	0	unknown	yes	

GitHub Location of Code:



Click Logo

Deployment:

AWS Infrastructure

1. Log in to the AWS portal

2. Creating a Data Stream

- a. After logging in to the AWS portal, search Kinesis and open the Kinesis Portal, now click on the Data Stream tab to create the Data Stream.
- b. This process is relatively straightforward; there are two capacity modes which are on-demand and provisioned, where on-demand helps us to automatically scale the data stream and provisioned mode uses an initial set of set resources.

Data stream capacity [Info](#)

Capacity mode

☒ **On-demand**
Use this mode when your data stream's throughput requirements are unpredictable and variable. With on-demand mode, your data stream's capacity scales automatically.

☐ **Provisioned**
Use provisioned mode when you can reliably estimate throughput requirements of your data stream. With provisioned mode, your data stream's capacity is fixed.

Total data stream capacity
By default, data streams with on-demand mode scale throughput automatically to accommodate traffic of up to 200 MiB per second and 200,000 records per second for the write capacity. If traffic exceeds capacity, your data stream will throttle.

Write capacity
Maximum
200 MiB/second and 200,000 records/second

Read capacity
Maximum (per consumer)
400 MiB/second
Up to 2 default consumers. Use Enhanced Fan-Out (EFO) for more consumers. EFO supports adding upto 20 consumers, each having a dedicated throughput.

8

Amazon Kinesis > Data streams > Create data stream

Create data stream [Info](#)

Data stream configuration

Data stream name

Acceptable characters are uppercase and lowercase letters, numbers, underscores, hyphens and periods.

Data stream capacity [Info](#)

Capacity mode

☒ **On-demand**
 Use this mode when your data stream's throughput requirements are unpredictable and variable. With on-demand mode, your data stream's capacity scales automatically.

☐ **Provisioned**
 Use provisioned mode when you can reliably estimate throughput requirements of your data stream. With provisioned mode, your data stream's capacity is fixed.

Total data stream capacity
 By default, data streams with on-demand mode scale throughput automatically to accommodate traffic of up to 200 MB per second and 200,000 records per second for the write capacity. If traffic exceeds capacity, your data stream will throttle.

Write capacity
 Maximum
 200 MB/second and 200,000 records/second

Read capacity
 Maximum (per consumer)
 400 MB/second
 Up to 2 default consumers. Use [Enhanced Fan-Out \(EFO\)](#) for more consumers. EFO supports adding up to 20 consumers, each having a dedicated throughput.

☒ On-demand mode has a pay-per-throughput pricing model. See [Kinesis pricing for on-demand mode](#).

Data stream settings
 You can edit the settings after the data stream has been created and is in the active status.

Setting	Value	Editable after creation
Capacity mode	On-demand	Yes
Data retention period	1 day	Yes
Server-side encryption	Disabled	Yes
Monitoring enhanced metrics	Disabled	Yes
Tags	-	Yes

[Cancel](#) [Create data stream](#)

3. Creating a Delivery Stream

- After Data Stream is created click on the Data Stream inside Kinesis to create the new Delivery Stream.
- There are two delivery streams to be created one for the real-time scenario where it is used for ElasticSearch, another one is used for the batch flow, and S3 is the Destination.

Delivery streams (2)

Delete

Create delivery stream

Find delivery streams

< 1 >

	Name	Status	Creatio...	Source	Data tr...	Destina...	Destination
	delivery-stream	Active	November...	data-strea...	processDa...	Amazon ...	ed-domain
	Kinesis-data-stream-to-S3	Active	December...	data-strea...	mykinesis...	Amazon S3	kinisestos3bucket007

- The configuration of the first delivery stream is as shown in the image
 1) Source would be the Data Stream we just created

- ** Use the following configuration while creating the services**

Cluster configuration

Security configuration

Cluster health

Instance health

Auto-Tune

Logs

Indices

Tags

Connections

VPC endpoints

Packages

Notifications

Cluster configuration

Edit

Data nodes

Availability Zones

1-AZ

Instance type

15.xsmall.search

Number of nodes

1

Storage type

EBS

EBS volume type

General Purpose (SSD) - gp3

EBS volume size

10 GB

Provisioned IOPS

3000 IOPS

Provisioned Throughput (MB/s)

125 MB/s

Dedicated master nodes

Enabled

No

Warm and cold data storage

Ultrawarm data nodes enabled

No

Network

Access

Public

Custom endpoint

Enabled

No

Snapshot

Frequency

Hourly

Start hour

00:00 UTC (default)

Advanced cluster settings

Allow references to indices inside the body of HTTP requests.

Yes

Fielddata cache allocation

20

Max clause count

1024

10

Security configuration

Edit

<p>Fine-grained access control</p> <p>Enabled</p> <p>Yes</p> <p>Master user type</p> <p>Internal user database</p>	<p>Authentication for OpenSearch Dashboards/Kibana</p> <p>SAML enabled</p> <p>No</p> <p>Cognito enabled</p> <p>No</p> <p>Region</p> <p>US East (N. Virginia)</p>	<p>Encryption</p> <p>Required HTTPS</p> <p>Yes</p> <p>Node-to-node encryption</p> <p>Yes</p> <p>Encryption at rest</p> <p>Yes</p> <p>AWS KMS key</p> <p>arn:aws:kms:us-east-1:541077676580:key/f6841690-a2be-4feb-ba68-c7d6375a558c</p>
--------------------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Access policy [Info](#)

Policy

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "*"
      },
      "Action": "ec:*",
      "Resource": "arn:aws:ec2:us-east-1:541077676580:domain/*ed-domain*"
    }
  ]
}

```

[Copy policy code](#)

Delivery Stream1 (Part A)

Source settings

Kinesis data stream

[data-stream](#)

Transform records

Edit

Configure Kinesis Data Firehose to transform your record data.

Transform source records with AWS Lambda [Info](#)

<p>Data transformation</p> <p>Enabled</p> <p>Description</p> <p>-</p> <p>Buffer size</p> <p>1 MIB</p>	<p>Lambda function</p> <p>processDataFirehose</p> <p>Lambda function version</p> <p>\$LATEST</p>	<p>Runtime</p> <p>nodejs18.x</p> <p>Timeout</p> <p>1 minute</p>
-------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------	-----------------------------------------------------------------

Destination settings [Info](#)

Edit

Specify the destination settings for your delivery stream.

OpenSearch Service destination

<p>Domain</p> <p>ed-domain</p> <p>Retry duration</p> <p>300 seconds</p>	<p>Index</p> <p>index1</p>	<p>Index rotation</p> <p>No rotation</p>
-----------------------------------------------------------------------------------------	----------------------------	------------------------------------------

Buffer hints

Buffer size

5 MIB



Buffer interval

300 seconds

Delivery Stream1 (Part B)


Backup settings [info](#) Edit

Enabling source record backup ensures that source records can be recovered if record processing transformation does not produce the desired results.

Backup mode	S3 backup bucket	S3 backup bucket prefix
Failed data only	mydbbackup202901	 ex/
Buffer size	Buffer interval	S3 backup bucket error output prefix
5 MB	60 seconds	 ex/
Encryption for data records	Compression for data records	
Not enabled	Not enabled	

Server-side encryption (SSE) [info](#) Edit

You can use [AWS Key Management Service \(KMS\)](#) to create and manage Customer Master Keys (CMKs) and to control the use of encryption across a wide range of [AWS](#) services in your applications.

 To enable SSE for the delivery stream, view the data stream selected above, and enable SSE on it.

Destination error logs [info](#) Edit

Choose Enabled if you want Kinesis Data Firehose to log record delivery errors to CloudWatch Logs.

Amazon CloudWatch error logging

Enabled

Permissions [info](#) Edit

Kinesis Data Firehose uses this IAM role for all the permissions that the delivery stream needs. To specify different roles for the different permissions, use the API or the CLI.

IAM role

[KinesisDataFirehoseServiceRole-delivery-stre-us-east-1-1669680908488](#)

Delivery Stream2 (Part A)

Source settings

Kinesis data stream
data-stream

Transform and convert records

Configure Kinesis Data Firehose to transform and convert your record data.

Transform source records with AWS Lambda

Data transformation

Enabled

Description

An Amazon Kinesis Firehose stream processor that accesses the records in the input and returns them with a processing status.

Buffer size

1 MB

Lambda function

mykinesisinput

Lambda function version

LATEST

Runtime

nodejs14.x

Timeout

12 minutes 10 seconds

Convert record format

Record format conversion

Not enabled

Destination settings

Specify the destination settings for your delivery stream.

Amazon S3 destination

S3 bucket

kinestestbucket007

S3 bucket error output prefix

-

Dynamic partitioning

Dynamic partitioning

Not enabled

Deaggregation delimiter

-

Multi record deaggregation

Not enabled

New line delimiter

Not enabled

Multi record deaggregation type

-

Inline parsing for JSON

Not enabled

S3 bucket prefix

-

Buffer hints

Buffer size

5 MB

Buffer interval

300 seconds

Compression and encryption

Compression for data records

Not enabled

Encryption for data records

Not enabled

Dynamic partitioning retry

Retry duration

-

Delivery Stream2 (Part B)

Backup settings

Enabling source record backup ensures that source records can be recovered if record processing transformation does not produce the desired results.

Source record backup in Amazon S3

Not enabled

Server-side encryption (SSE)

You can use AWS Key Management Service (KMS) to create and manage Customer Master Keys (CMKs) and to control the use of encryption across a wide range of AWS services in your applications.

To enable SSE for the delivery stream, view the data stream selected above, and enable SSE on it.

Destination error logs

Choose Enabled if you want Kinesis Data Firehose to log record delivery errors to CloudWatch Logs.

Amazon CloudWatch error logging

Enabled

Permissions

Kinesis Data Firehose uses this IAM role for all the permissions that the delivery stream needs. To specify different roles for the different permissions, use the API or the CLI.

IAM role

KinesisDataFirehoseServiceRole-Kinesis-data-us-east-1-1669947217318

Tags (0)

You can add tags to organize your AWS resources, track costs, and control access.

Key

Value

No tags

No tags associated with this stream.

4. Creating a Lambda Function

- After Delivery Stream is created, search for Lambda, click on create a new function and then create and deploy the function.

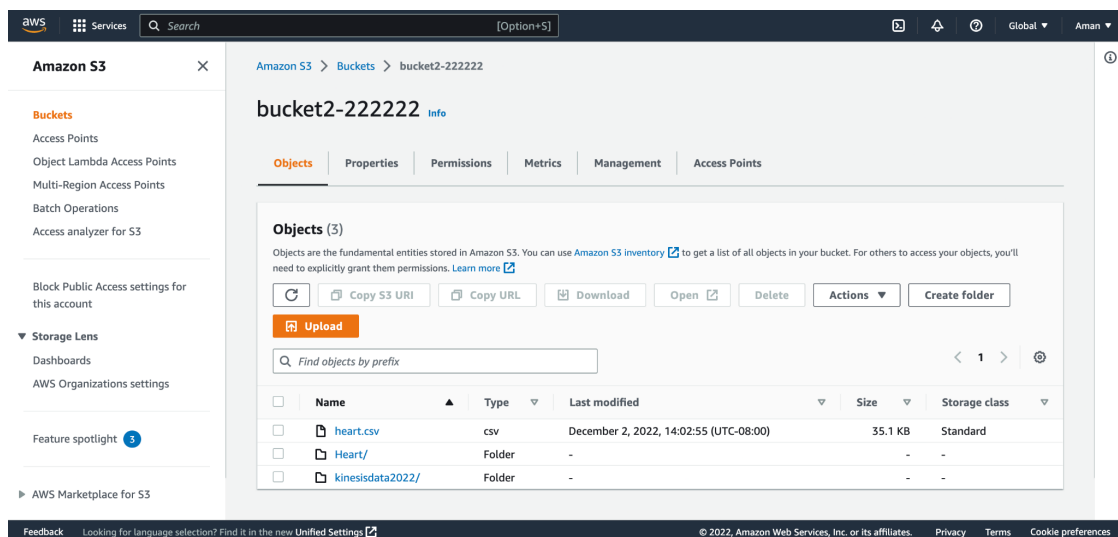
```

4 let success = 0; // Number of valid entries found
5 let failure = 0; // Number of invalid entries found
6 /* Process the list of records and transform them */
7 const output = event.records.map((record) => {
8   // Kinesis data is base64 encoded so decode here
9   console.log(record.recordId);
10  const payload = (Buffer.from(record.data, 'base64')).toString('ascii');
11  console.log('Decoded payload:', payload);
12  // Split the data into it's fields so we can refer to them by index
13  const match = payload.split('|');
14
15  if (match) {
16    /* Prepare JSON version from Syslog log data */
17    const result = {
18      // build all fields from array
19      Age: match[0],
20      Job: match[1],
21      Marital: match[2],
22      Education: match[3],
23      Default: match[4],
24      Balance: parseInt(match[5]),
25      Housing: match[6],
26      Loan: match[7],
27      Contact: match[8],
28      Days: parseInt(match[9]),
29      Month: match[10],
30      Duration: match[11],
31      Campaign: match[12],
32      Pdays: parseInt(match[13]),
33      Previous: parseInt(match[14]),
34      Poutcome: match[15],
35      Deposit: match[16]
36    };
37    console.log('Result:', payload);
38    success++;
39    return {
40      recordId: record.recordId,
41      result: 'OK',
42      data: (Buffer.from(JSON.stringify(result))).toString('base64'),
43    };
44  } else {
45    /* Failed event, notify the error and leave the record intact */
46    failure++;
47    return {
48      recordId: record.recordId,
49      result: 'ProcessingFailed',
50      data: record.data,
51    };
52  }
53 });
54 console.log('Processing completed. Successful records ${success}, Failed records ${failure}.');
55 callBack(null, { records: output });

```

5. Creating S3 storage

- After Lambda Function is Created, search for Lambda, click on create a new function and then create and deploy the function.



6. Finish creating the delivery stream

- After S3 and Lambda are created we can finish creating the delivery stream

7. Create a Glue Crawler

- Create a new Glue Crawler with the following configuration which would create a GLUE database and tables/

Glue Crawler

Crawler properties			
Name extractJson	IAM role AWSGlueServiceRole-GlueRole ↗	Database bankdata	State READY
Description -	Security configuration -	Lake Formation configuration -	Table prefix -
Maximum table threshold -			
▼ Advanced settings			
Create single schema for each S3 path False	Inherit schema from table False	Table level -	Schema updates in the data store Update the table definition in the data catalog for all data stores except S3. For tables that map to S3 data, add new columns only.
Object deletion in the data store Mark the table as deprecated in the data catalog.	Repeat crawls of S3 data stores Crawl all folders again with every subsequent crawl.		

Glue Database

Data Catalog

Databases New

Tables New

Stream schema registries

Schemas

Connections New

Crawlers New

Classifiers New

Catalog settings

▼ Data Integration and ETL

AWS Glue Studio

Jobs New

Interactive Sessions

Notebooks New

Data classification tools

Sensitive data detection 🔗

Record Matching

Triggers

Workflows

Blueprints

Security configurations

► Legacy pages

What's New New

AWS Glue > Tables > View table details

2022

Page last updated: December 2, 2022 at 22:09:05 (UTC)

Version 1 (Current version) ▼

🔄

Actions ▼

Table overview | Data quality

Table details | Advanced properties

Name 2022	Description -	Database bankiddata	Classification joon
Location s3://kmsiextos3bucket007/2022/	Connection -	Deprecated -	Last updated December 2, 2022 at 22:09:05
Input format org.apache.hadoop.mapred.TextInputFormat	Output format org.apache.hadoop.hive.q1.is.HiveIgnoreKeyTextOutputFormat	Serde serialization lib org.openx.data.joinserde.JsonSerDe	

Schema | Partitions | Indexes

Schema (20)

View and manage the table schema.

Edit schema as JSON Edit schema

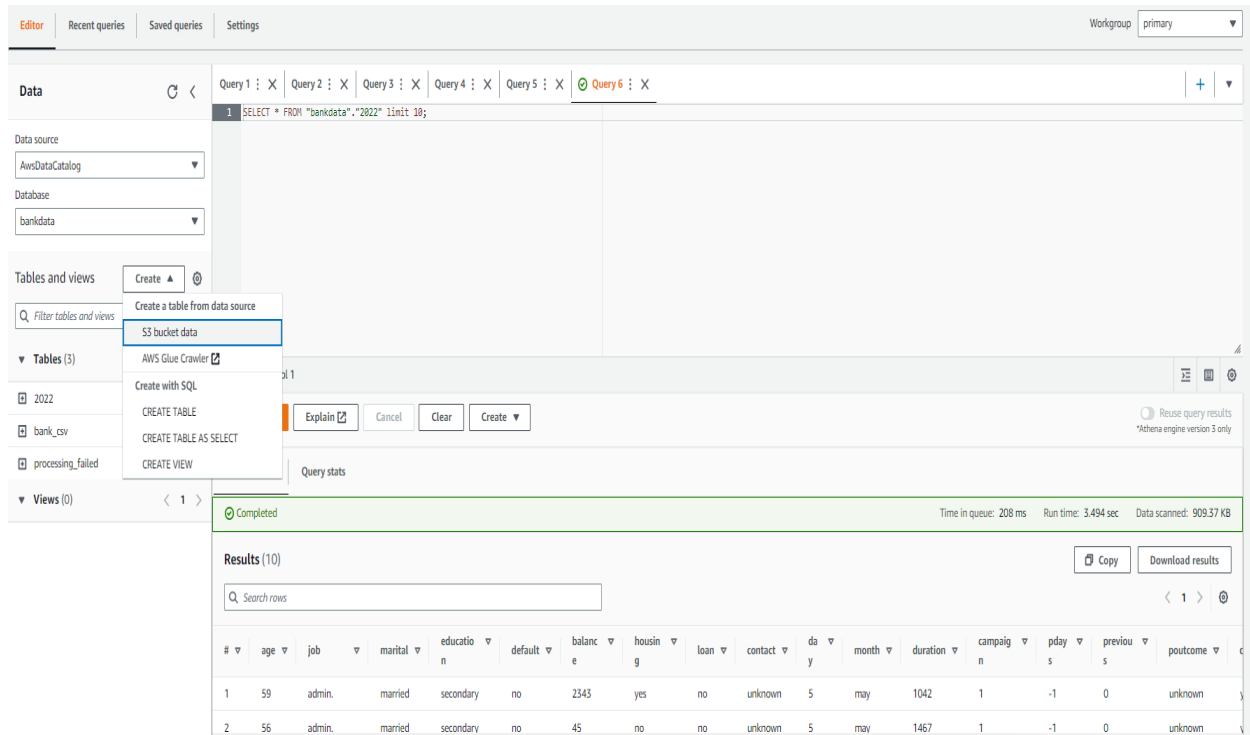
< 1 > ⌕

#	Column name	Data type	Partition key	Comment
1	age	string	-	-
2	job	string	-	-

8. Setting up Amazon Athena

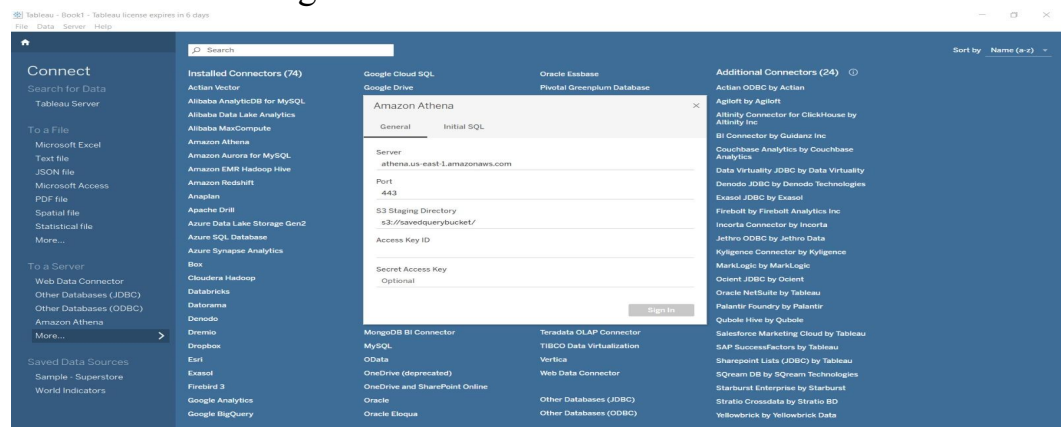
- After Creating AWS Glue Crawler, we need to search for Athena

- After a portal similar to the below image is opened we need to click on AWS Glue Crawler and select the already created Glue crawler.



9. Setting up Tableau

- After completing all the above steps, we need to install tableau and use the below configuration to connect Tableau to the Glue tables.



Steps to Run the Application

- 1) Download AWS CLI from <https://AWS.amazon.com/cli/>
- 2) Create an IAM role and give administrator access (for development) or give suitable access and note down the Access Key ID and the Access key
- 3) Open the Terminal(Command Prompt) and navigate to the project directory.
- 4) Type AWS configure and add the previously noted AWS Key ID and AWS Key, region, and output format.
- 5) Now set up the infrastructure in AWS, as explained above in the deployment stage.
- 6) After AWS has been set up, run the python file, which has the producer code, and will start adding the data into the AWS Data Streams and eventually in the AWS Firehose.
* Run **python <Python File name>.py**
* Generator Code and Amazon Lambda code can be found in the mentioned GitHub repository.
- 7) To view the Dashboard, we have to click on the Kibana URL inside the Domain, which we will find inside the created Delivery Stream

Test Results

Kibana Analysis

Below are the results of the live data stream analysis in the form of Dashboard

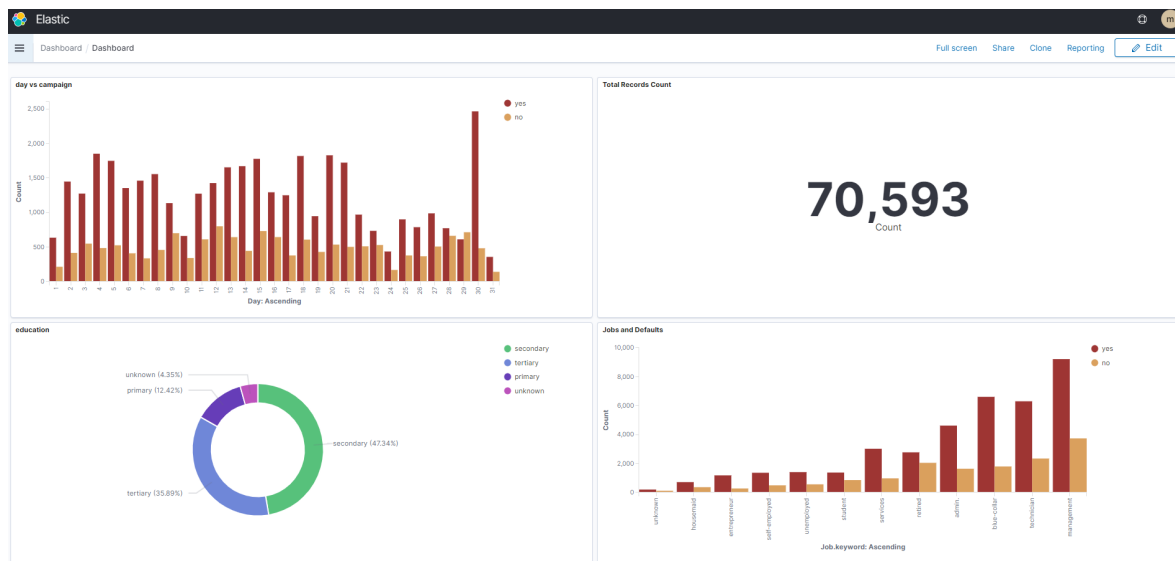
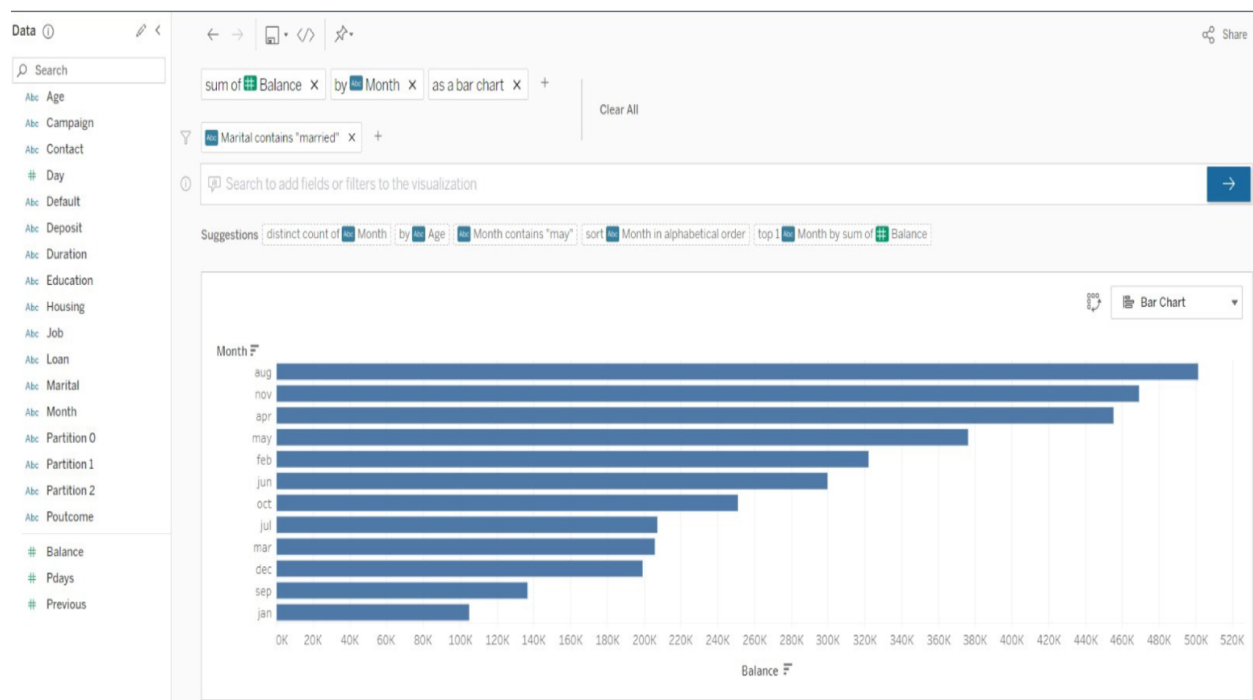
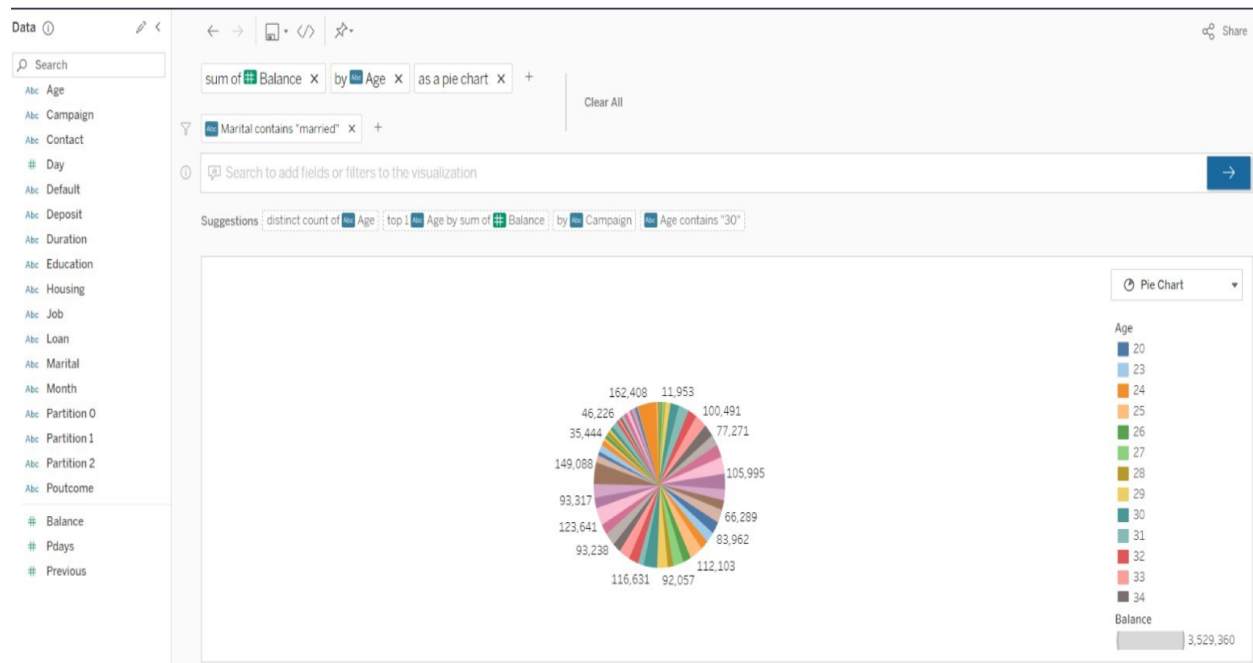


Tableau Analysis

Below are the results for the raw data analysis



References

<https://aws.amazon.com/kinesis/>

<https://aws.amazon.com/kinesis/data-firehose/>

<https://aws.amazon.com/opensearch-service/>

<https://aws.amazon.com/s3/>

<https://aws.amazon.com/glue/>

<https://aws.amazon.com/lambda/>

<https://aws.amazon.com/athena/>

<https://www.tableau.com/>